



Container security: best practices & roadmap

October 2019

Google Cloud

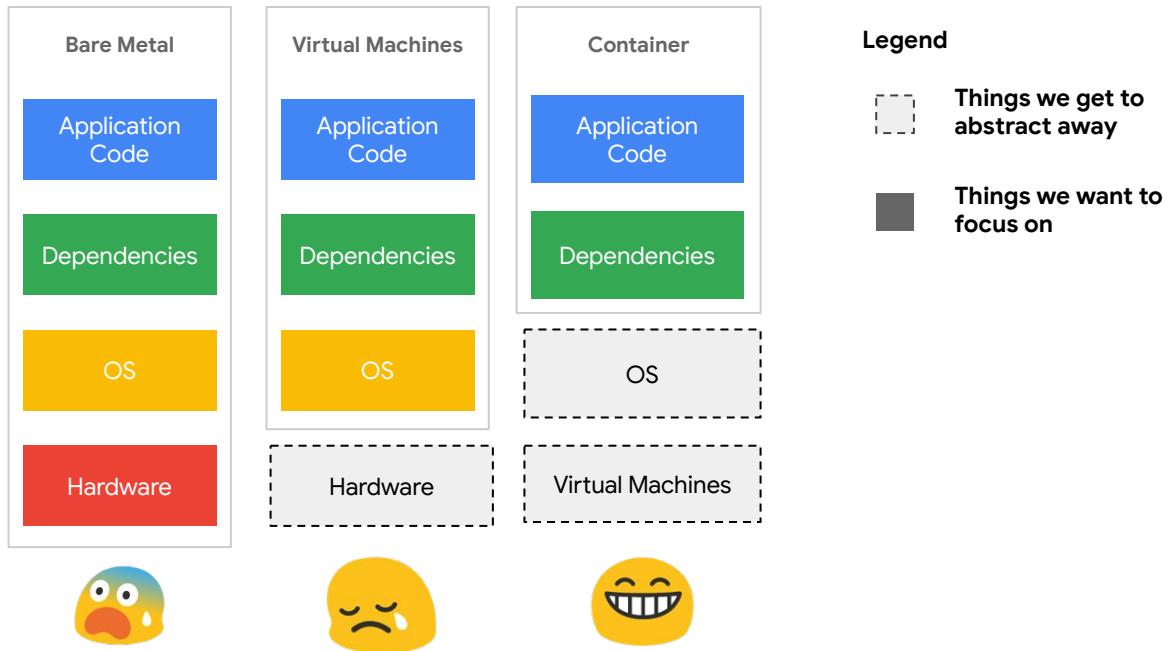
Container security 101



First, what's a container?

Containers are all the rage these days, getting us close to the dream of write once, run anywhere.

They package and isolate applications and dependencies.



What is Kubernetes?

- A portable, open-source, **container-centric** management platform
- Built-in primitives for **deployments, rolling upgrades, scaling, monitoring, and more**
- Inspired by **Google's internal systems**
- Get true **workload portability** and increased **infrastructure efficiency**



Containers come with security advantages inherent in their architecture



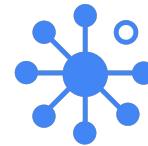
Containers are short lived and frequently re-deployed; **you can constantly be patching.**



Containers are intentionally immutable; a modified container is a **built-in security alert.**



Good security defaults are one line changes; **setting secure configurations is easy.**



With isolation technologies, **you can increase security without adding resources.**

Container security threats & risks

INFRASTRUCTURE SECURITY

- Privilege escalation
- Credential compromise
- Kubernetes API compromise
- Over-privileged users

SOFTWARE SUPPLY CHAIN

- Unpatched vulnerability
- Supply chain vulnerability
- Zero day exploit on common library

RUNTIME SECURITY

- DDoS
- Node compromise and exploit
- Container escape
- Flood event pipeline

What is container security?

INFRASTRUCTURE SECURITY

Is my infrastructure **secure for developing containers?**

- How can I use Kubernetes security features to protect my identities, secrets, and network?
- How can I use native GCP functionality, like IAM, audit logging, and networking?

SOFTWARE SUPPLY CHAIN

Is my container image **secure to build and deploy?**

- How can I make sure my container images are vulnerability-free?
- How can I make sure the images I built aren't modified before they are deployed?

RUNTIME SECURITY

Is my container **secure to run?**

- How can I identify a container acting maliciously in production?
- How can I take action to protect and isolate my workload?
- How can I securely scale my containers deployment?

Kubernetes security features



Namespaces	RBAC	NetworkPolicy
Segment cluster resources	Control access to cluster resources	Limit pod to pod traffic
Pod Security Policy	Secrets	Audit Logging
Limit what permissions pods have	Provide pods access to application secrets	Capture operations to the Kubernetes API server



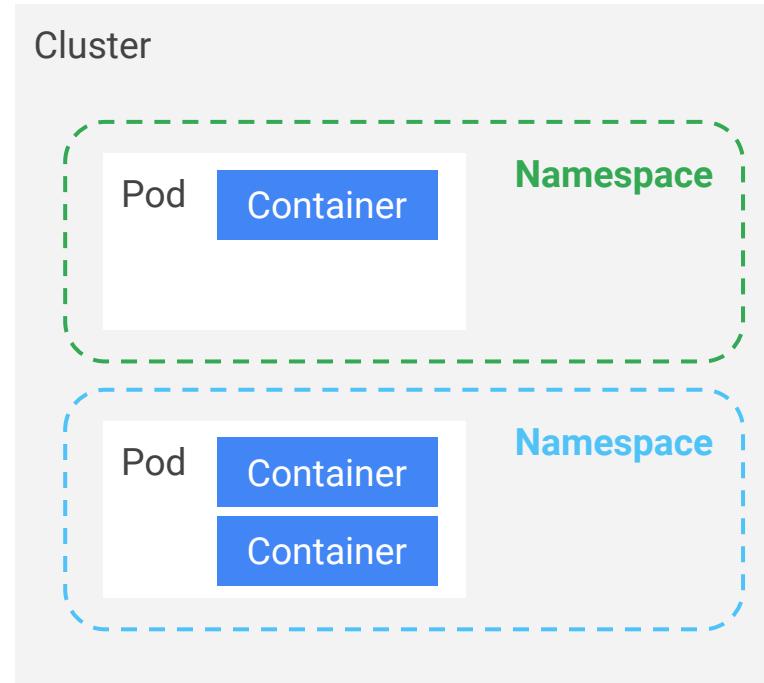
K8S: Kubernetes namespaces

Problem: I have too much stuff!

- Name collisions in the API
- Poor isolation between users
- Don't want to expose sensitive data

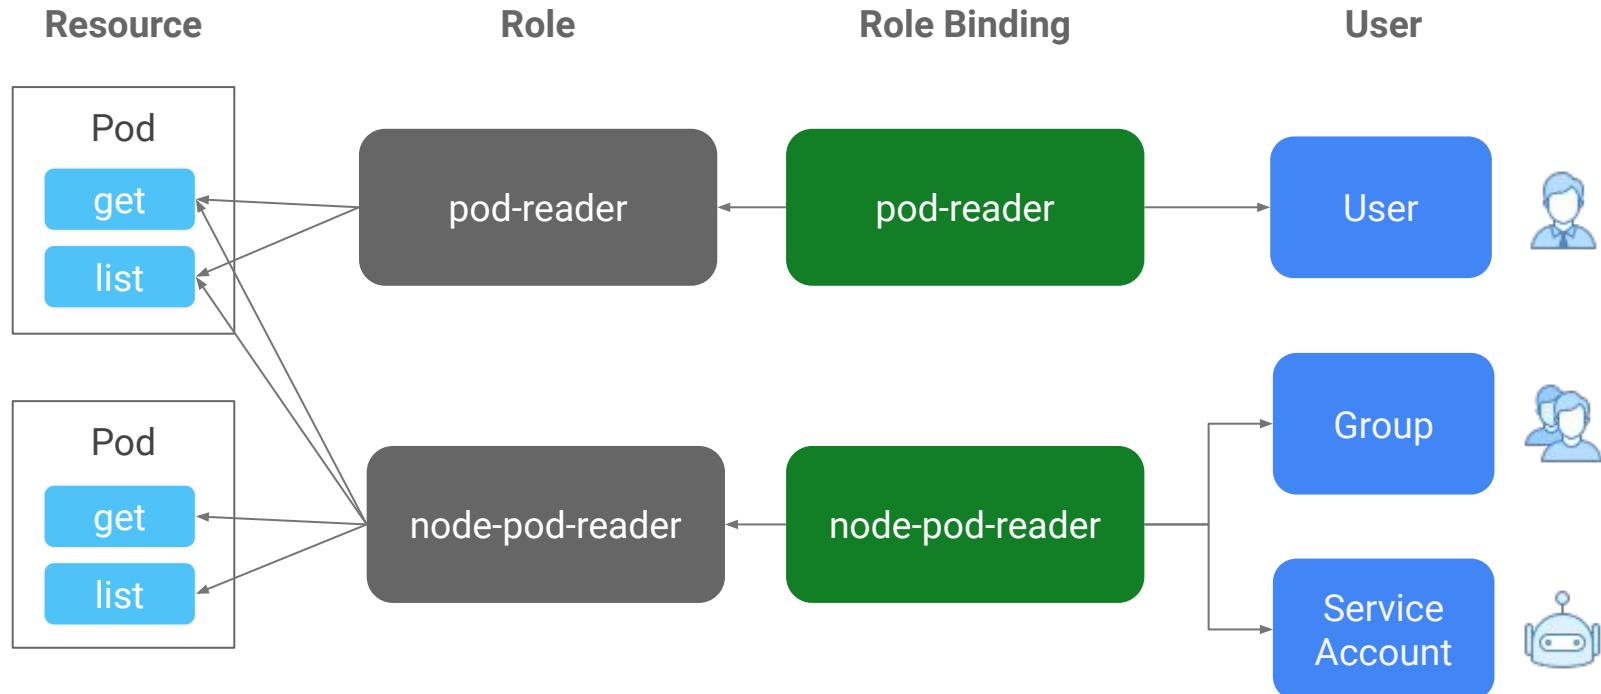
Solution: Slice up the cluster

- Create new namespaces as needed
 - per-user, per-app, per-department, etc.
- Namespaces are just another API object
- Provide each user community with its own:
 - Resources, with consumption limits
 - Policies, with delegated management





K8S: RBAC

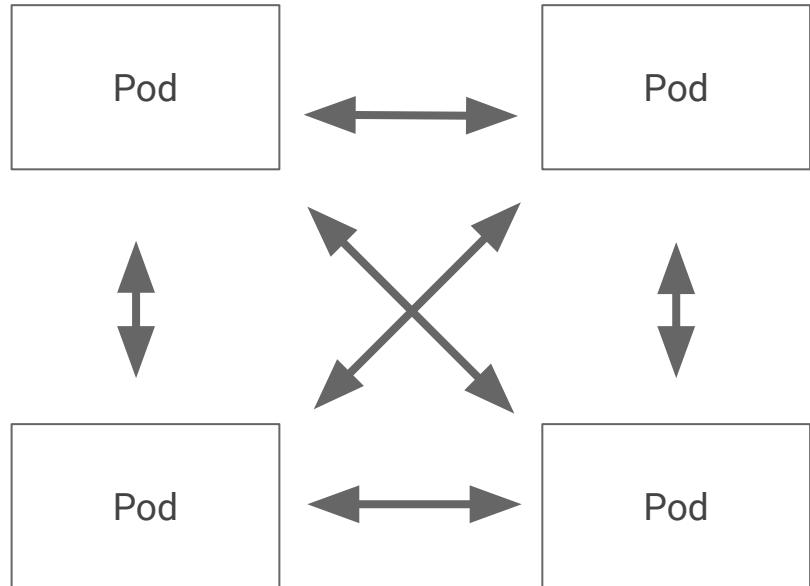




K8S: NetworkPolicy

How does K8S networking work?

- IPs are per pod, **scoped to cluster**
- **Pods can reach each other directly**, without NAT, even across nodes





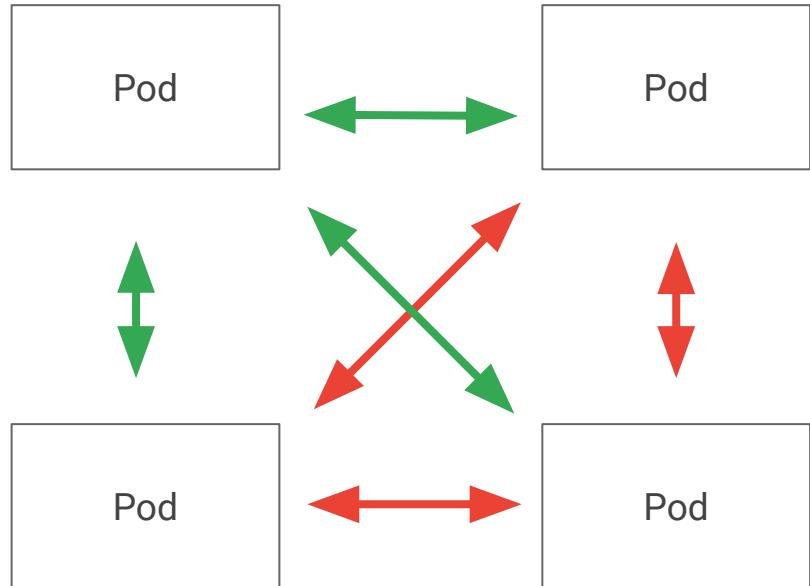
K8S: NetworkPolicy

How does K8S networking work?

- IPs are per pod, **scoped to cluster**
- Pods can reach each other **directly**, without NAT, even across nodes

With NetworkPolicy:

Restrict pod-to-pod traffic



Best practices and current features



Google's container security philosophy

- Prioritize security by default**
We believe in giving users the power of Kubernetes without making them security experts
- Integrate the best of Google Cloud Platform security**
Integrations with IAM, Audit Logging, VPCs, encryption by default, node auto-upgrades, and the Cloud Security Command Center
- Take advantage of a strong security culture**
From our dedicated infra security team to security design requirements, security is an integral part of Google's culture, products, and operations

Container security products & features

INFRASTRUCTURE SECURITY	SOFTWARE SUPPLY CHAIN	RUNTIME SECURITY
Best practices by default RBAC, disabled Kubernetes dashboard, and hardening guide	Workload identity ^{BETA} Google IAM service accounts for pods to authenticate outside the cluster	Container Registry Vulnerability Scanning Automatic scans of images and packages for known CVEs
Minimal OS Hardened, purpose-built host OS based on Chromium	Shielded GKE nodes ^{BETA} Verify boot integrity and secure kubelet credentials using vTPMs	Binary Authorization Verification of provenance and blocking deployments that do not meet necessary requirements
Node auto-upgrade Automatic Kubernetes version upgrades for nodes, including security patches	Customer-managed encryption keys ^{BETA} Encryption of persistent disks with a key in Cloud KMS	Managed base images Base images for common distributions patched for security vulnerabilities automatically by Google
Private clusters & authorized networks Private IPs for your nodes, and restricted IP access to masters	Application-layer secrets encryption ^{BETA} Manage encryption of secrets in Kubernetes using a key in Cloud KMS	Pod Security Policy ^{BETA} Pod capability constraints, including Linux constructs like seccomp
Managed SSL certs ^{BETA} HTTPs load balancers with automatically configured SSL certificates		





GKE: Node auto-upgrade

Node upgrades help you keep your nodes up to date with the latest stable version of Kubernetes

- When a security patch is rolled out to nodes, GKE automatically patches your node and keeps it up to date
 - Security patch is rolled into the first possible release after public disclosure
 - Auto-upgrade is typically rolled out 2-4 weeks after update, depending on the severity of the vulnerability
- Available for both COS and Ubuntu nodes
- You can also schedule a maintenance window to control when maintenance happens



GKE: Minimal OS

Container-optimized OS (COS) based on Chromium OS, and maintained by Google

- **Built from source:** Since COS is based on Chromium OS, Google maintains all components and is able to rebuild from source if a new vulnerability is discovered and needs to be patched
- **Smaller attack surface:** Container-Optimized OS is purpose-built to run containers, has a smaller footprint, reducing your instance's potential attack surface
- **Locked-down by default:** Firewall restricts all TCP/UDP except SSH on port 22, and prevents kernel modules. Root file system is mounted read-only
- **Automatic Updates:** COS instances automatically download weekly updates in the background; only a reboot is necessary to use the latest updates. Google provides patches and maintenance

GKE leverages Google Cloud security features



Identity & Access Management

Manage your users and permissions at the project level

This doesn't stop you from using Kubernetes RBAC - but lets you more easily manage this across projects



Cloud Audit Logging

Review, monitor and alert on audit logs centrally

This makes Kubernetes audit logs easier to consume along with all your other audit logs



Virtual Private Cloud

Connect your cloud resources to each other on a private network

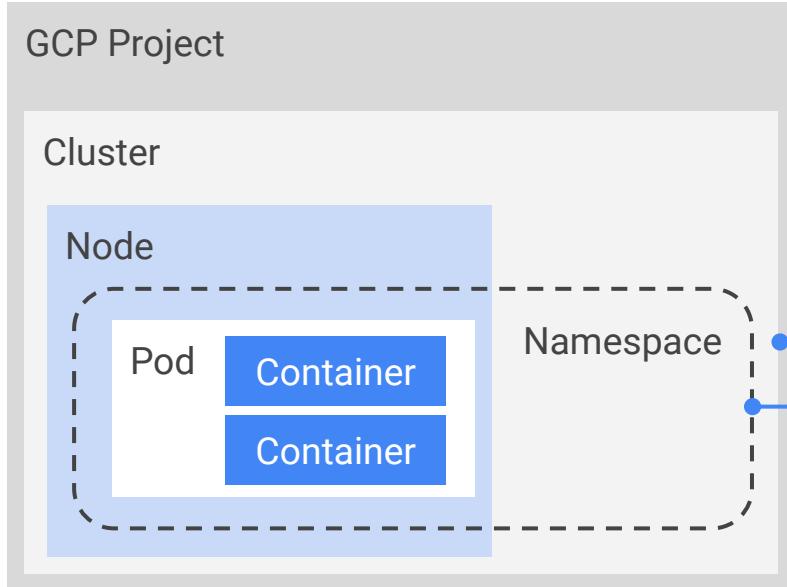
This lets you plug in Kubernetes into your broader network, the way you want to



Google Cloud



GKE: Using IAM and RBAC



Use IAM at the project level

Set roles for

- Cluster Admin: manage clusters
- Container Developer: API access within clusters

Use RBAC at the cluster and namespace level

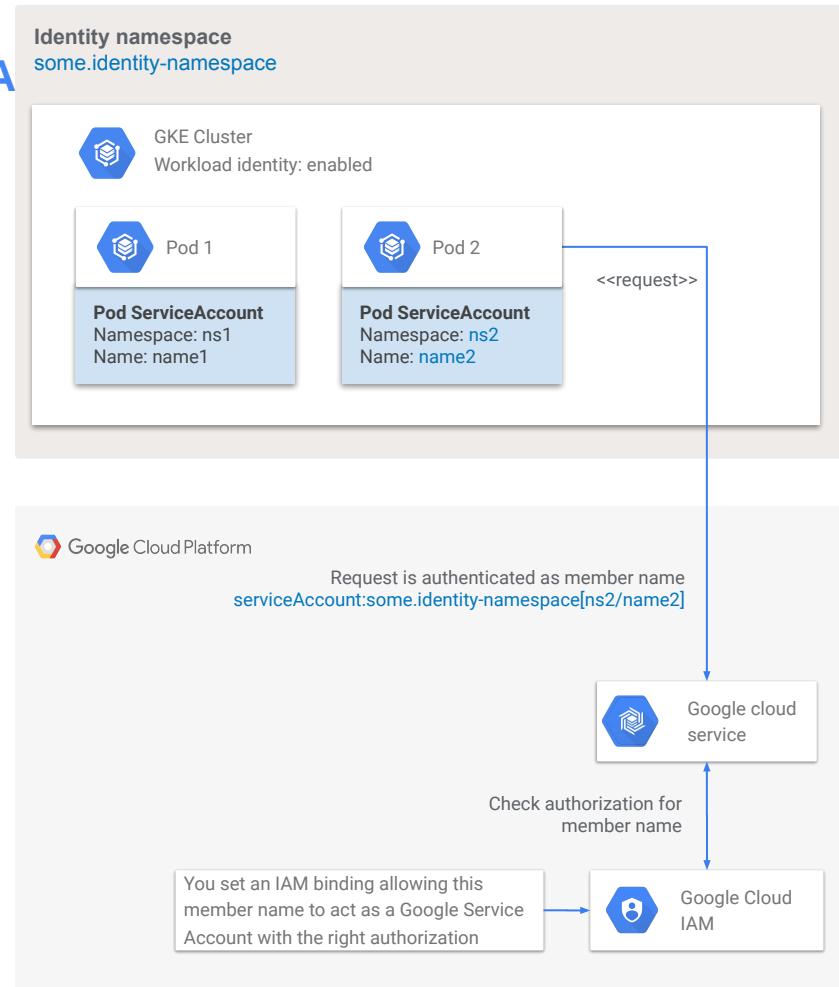
Set permissions on individual clusters and namespaces



GKE: Workload identity BETA

Access Google Cloud services from Google Kubernetes Engine without managing any keys or secrets at all

- New IAM primitive, ‘identity namespace,’ enables Google cloud platform to authenticate Kubernetes Service Accounts directly
 - Use IAM to grant the desired authorization for a kubernetes service account
 - Google manages the identity namespace credentials; no rotation burden
 - Each workload only has minimal, necessary permissions
- Easy to assign identity and prove it to Google services





GKE: Shielded GKE nodes BETA

- **Crypto anchors with vTPMs** offers isolation of secrets bound to a TPM, and exfiltration resistance
- How does it work?
 - Node Identity will be attested to ensure no arbitrary machine can join a cluster as a node
 - Any secret can be bound to a node (vTPM) and will not be globally shared



GKE: Application layer secrets encryption BETA

Google Cloud Key Management Service



Google Kubernetes Engine

Master



Node



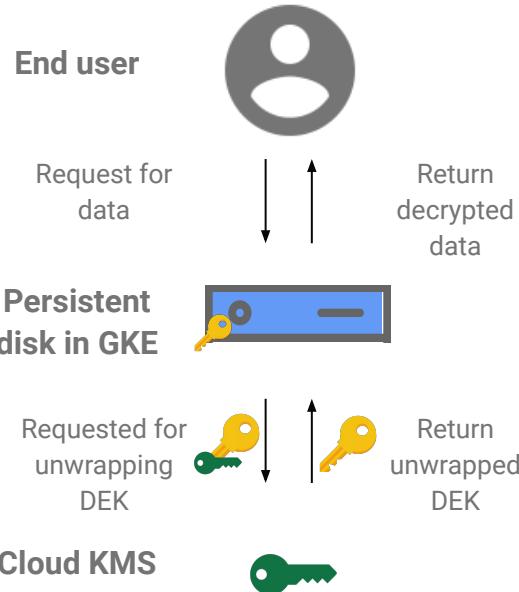
- Data in etcd, such as secrets, are encrypted locally with a data encryption key (DEK)
- The data encryption key is also stored in etcd, but encrypted with a key encryption key (KEK) in Cloud KMS, via a KMS plugin
- Each encryption generates a new DEK, rather than rotating the DEK



Google Cloud



GKE: Customer-managed encryption keys BETA



- Encryption key per cluster, for the GCE disks in that cluster
- You specify the Cloud KMS key you want to use
 - Must be in the same region
 - Must give permissions on the key to the GKE service account
- Key is specified as part of storage class
- Limitations
 - Only for dynamically attached persistent disks
 - Requires self-installation of CSI driver



Managed base images

- Managed base images are patched for security vulnerabilities automatically by Google with the most recently available patches from upstream
- Most common distributions: Available for Debian, Ubuntu, and CentOS, on GCP Marketplace
- Scanned and patched: Scanned for vulnerabilities using Container Registry Vulnerability Scanning. Where a patch for these exists upstream, it is applied.
- Used by Google: Google maintains several base image for building its own applications, including Google Cloud services like Google App Engine



GCR: Vulnerability Scanning

- Scans all images in your private Google Container Registry for known Common Vulnerabilities and Exposures (CVEs)
- Examines images and packages installed in images
- Works for: Debian, Ubuntu and Alpine images
- Images are scanned when:
 - An image is added to the registry
 - There is an update to the vulnerability database



Binary Authorization

Run only what you trust. Enforce that only images with required signatures, or from whitelisted repositories, can be deployed to production.

- Customer defines signatures requirements
- Images are signed throughout CI/CD process by build/test tools
- Image deployments are blocked/allowed at GKE control plane based on policy

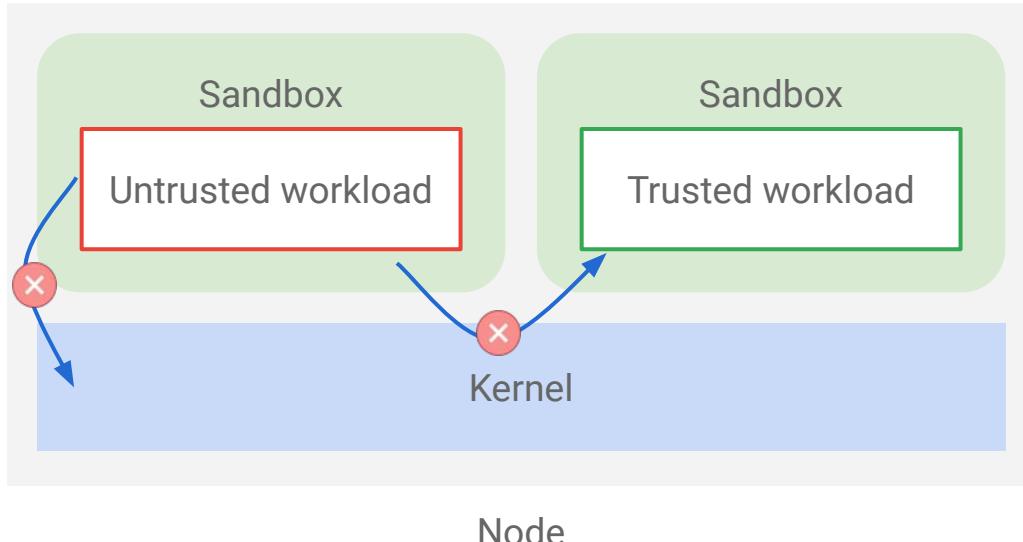


GKE: GKE Sandbox^{BETA}

Run **trusted and untrusted** workloads
on the same node

Rather than achieving isolation via
separate VMs, you can run workloads
of different trust levels on the same
node

Performance improvements from not
having to allocate a new cluster to
achieve isolation



Who protects what?

Hardening the control plane is **Google's responsibility**

Google does most of the **hard work** to protect nodes - but it's the **user's responsibility to upgrade** to reap the benefits

Protecting workloads is still the **user's responsibility**

Responding to attacks on **Google's infrastructure** is **Google's responsibility**, whereas responding to attacks on **user workloads** is the **user's responsibility**

**Upcoming
products
and
features
Under NDA**



Container security: What's next?

INFRASTRUCTURE SECURITY

Tests in Security Health Analytics^{ALPHA}
Automated checks for configurations following the GKE hardening guide

GKE CIS Benchmarks
Updated CIS benchmarks for Kubernetes, as well as new CIS benchmarks for GKE specifically

SOFTWARE SUPPLY CHAIN

Vulnerability deployment policy^{ALPHA}
Enforce deployment policies based on the vulnerabilities and base images used

Deployment History^{EAP}
Verify if your environment is affected by a newly discovered vulnerability

RUNTIME SECURITY

Container Threat Detection^{ALPHA}
Monitor, detect, and react to common container attacks

Container security

Infrastructure security

Tests in Security Health
Analytics

Automated checks for
configurations following the
GKE hardening guide



CSCC: Security Health Analytics ALPHA

Container Scanner

Category	Finding description
IP_ALIAS_DISABLED	Indicates that a GKE Cluster was created with Alias IP ranges enabled.
LEGACY_AUTHORIZATION_ENABLED	Indicates that Legacy Authorization is enabled on GKE Clusters.
MASTERAUTHORIZED_NETWORKS_DISABLED	Indicates that Master authorized networks is not enabled on GKE Clusters.
MONITORING_DISABLED	Indicates that Stackdriver Monitoring is disabled on GKE Clusters.
NETWORK_POLICY_DISABLED	Indicates that Network policy is disabled on GKE Clusters.
POD_SECURITY_POLICY_DISABLED	Indicates that PodSecurityPolicy is disabled on a GKE Cluster.
PRIVATE_CLUSTER_DISABLED	Indicates that a GKE Cluster has a Private cluster disabled.
WEB_UI_ENABLED	Indicates that the GKE web UI (dashboard) is enabled.

- Checks for common misconfigurations in GCP
 - Such as those in the GKE hardening guide
- Results shown in the Cloud Security Command Center

Container security

Software supply chain

Vulnerability deployment policy

Enforce deployment policies based on the vulnerabilities and base images used

Deployment History

Verify if your environment is affected by a newly discovered vulnerability

Signing keys in Cloud KMS

Manage signing keys in Binary Authorization in Cloud KMS



GCR: Deployment History EAP

- Works for all images stored in GCR and images deployed to GCP from other registries
- Automatically detects when container images are running in GKE or App Engine Flex
- Makes information available via the Container Analysis API about what's running and has been running in a project
- Enables identification of known vulnerabilities that apply to running containers

Questions

- How do you know what is currently deployed in your environment? And associated vulns?

Container security

Runtime security

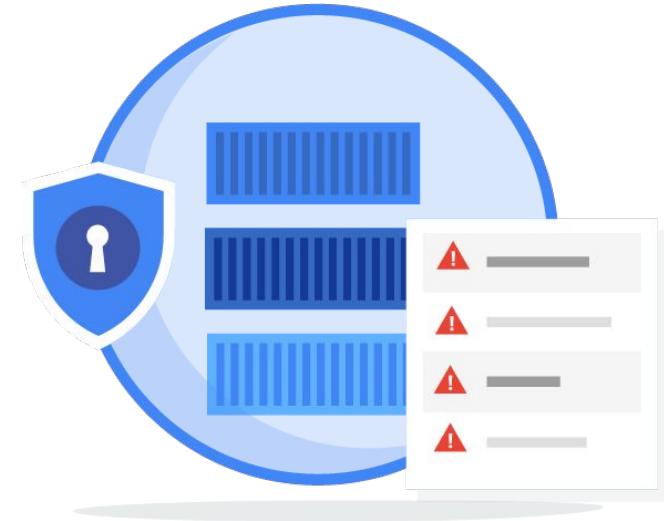
Container Threat Detection

Monitor, detect, and react
to common container
attacks

Container Threat Detection^{Alpha}

Container Threat Detection will offer detection capabilities
that cover the most common threats

- Pre-built detectors for the top attacks
 - Suspicious binary
 - Suspicious library
 - Reverse shell
- Integrations with
 - Cloud Security Command Center for surfacing security findings
 - GKE for one-click deployment
 - Stackdriver for events
- Kernel integrations only available in COS



Learn more

cloud.google.com/containers/security

g.co/gke/security

g.co/gke/hardening

g.co/gke/securitybulletins

